LiteRAG Deployment Guide

1. Prerequisites

Java 17 or later, Node.js 18 or later, Docker (optional).

Qdrant runs as a single binary on port 6333 (or via docker compose).

2. Start the backend

Run: java -jar literag-server.jar --server.port=8080

The metadata database is an embedded SQLite file at ./data/db.

3. Start the frontend

Run: npm install && npm run dev inside the web directory.

The dev server proxies /api and /mcp to port 8080.

4. Configure an embedding model

Open the Models page, add an OpenAI compatible endpoint
such as BAAI/bge-m3, test the connection and set it as default.

5. Connect an MCP client

Point the client at http://localhost:8080/mcp using the
Streamable HTTP transport, then call search_knowledge.